

Table 1. Full HarmBench robust-refusal results for Llama-3.2-1B-Instruct. Direct requests and each of the five frozen human-jailbreak templates contain the same 240 behavior identities.

Model	Template	Method	ASR (%) ↓	Safe-concept relevance (0–2) ↑	Instruction relevance (0–2) ↑	Fluency (0–2) ↑	Overall steering (0–2) ↑
Llama-3.2-1B-Instruct	Direct	Original	10.00	1.91	1.73	1.85	1.62
		A-LQR	9.58	1.95	<u>1.51</u>	<u>1.89</u>	<u>1.46</u>
		H_∞ (ours)	2.92	1.96	1.29	1.74	1.22
	John persona	Original	5.83	1.97	0.68	1.60	0.65
		A-LQR	6.25	1.97	<u>0.62</u>	<u>1.52</u>	<u>0.55</u>
		H_∞ (ours)	1.25	1.98	0.39	1.19	0.34
	DNE nonresponse	Original	6.25	1.99	0.63	1.59	0.64
		A-LQR	4.17	<u>1.99</u>	<u>0.57</u>	<u>1.53</u>	<u>0.56</u>
		H_∞ (ours)	1.67	1.98	0.38	1.18	0.30
	Jailbreak Bot	Original	4.58	1.96	0.64	1.62	0.63
		A-LQR	3.33	1.96	<u>0.56</u>	<u>1.53</u>	<u>0.53</u>
		H_∞ (ours)	0.42	1.98	0.38	1.17	0.30
	YOJA/Nona roleplay	Original	8.33	1.97	0.63	1.56	0.62
		A-LQR	8.75	1.97	<u>0.55</u>	<u>1.50</u>	<u>0.51</u>
		H_∞ (ours)	0.83	1.98	0.37	1.15	0.29
	APM programmer	Original	2.92	1.94	0.64	1.56	0.64
		A-LQR	5.00	1.97	<u>0.57</u>	<u>1.51</u>	<u>0.53</u>
		H_∞ (ours)	0.42	1.98	0.37	1.12	0.30