

Table 1. HarmBench attack success rate by model scale, reported separately for direct requests and each frozen human-jailbreak template. Human average is the equal-weight mean over the five templates; TBD denotes an incomplete model evaluation.

Method		Attack success rate (%) ↓						
		Direct	John	DNE	Jailbreak Bot	YOJA/Nona	APM	Human jailbreak average
Llama-3.2-1B-Instruct	Original	10.00	5.83	6.25	4.58	8.33	2.92	5.58
	A-LQR	9.58	6.25	4.17	3.33	8.75	5.00	5.50
	H_∞ (ours)	2.92	1.25	1.67	0.42	0.83	0.42	0.92
Llama-3.2-3B-Instruct	Original	16.67	11.67	10.83	14.17	11.25	10.83	11.75
	A-LQR	20.00	11.67	10.42	12.92	11.67	11.25	11.58
	H_∞ (ours)	8.33	4.17	3.75	3.75	4.17	2.50	3.67
Llama-3.1-8B-Instruct	Original	15.83	23.33	14.17	9.17	15.00	11.67	14.67
	A-LQR	15.83	15.00	11.67	7.50	9.17	10.83	10.83
	H_∞ (ours)	3.33	1.67	0.00	0.00	0.00	0.83	0.50