

Table 1. Full L-CiteEval length-transfer results for Qwen2.5-3B-Instruct. Each context length uses the same 40 question identities. Higher is better for every column.

Context	Method	Answer recall (%) $\uparrow$	Citation F1 (%) $\uparrow$	Concept relevance (0–2) $\uparrow$	Instruction relevance (0–2) $\uparrow$	Fluency (0–2) $\uparrow$
8K	Original	66.2	7.8	0.00	1.35	1.50
	S-PID	<u>72.1</u>	<u>7.2</u>	0.00	<u>1.60</u>	1.50
	A-LQR	70.1	5.6	<u>0.05</u>	1.48	1.48
	H_∞ (ours)	66.4	6.4	0.00	1.55	1.55
16K	Original	65.2	5.8	0.20	1.40	1.52
	S-PID	61.9	7.2	0.00	1.38	1.48
	A-LQR	64.8	3.0	0.00	<u>1.43</u>	<u>1.57</u>
	H_∞ (ours)	65.2	7.8	0.15	1.40	1.43