

Table 1. Summary L-CiteEval length-transfer results for Qwen2.5-3B-Instruct, macro-averaged equally across the matched 8K and 16K conditions. Higher is better for every column.

Qwen2.5-3B-Instruct	Method	Answer recall (%) $\uparrow$	Citation F1 (%) $\uparrow$	Concept relevance (0–2) $\uparrow$	Instruction relevance (0–2) $\uparrow$	Fluency (0–2) $\uparrow$
	Original	65.7	6.8	0.10	1.38	1.51
	S-PID	67.0	<u>7.2</u>	0.00	<u>1.49</u>	1.49
	A-LQR	<u>67.4</u>	4.3	0.03	1.45	<u>1.52</u>
	H_∞ (ours)	65.8	7.1	0.07	1.48	1.49