


Control-theoretic steering of LLMs: truthfulness and toxicity summary

A. Truthfulness across model scales



B. Toxicity steering (DistilGPT-2)

